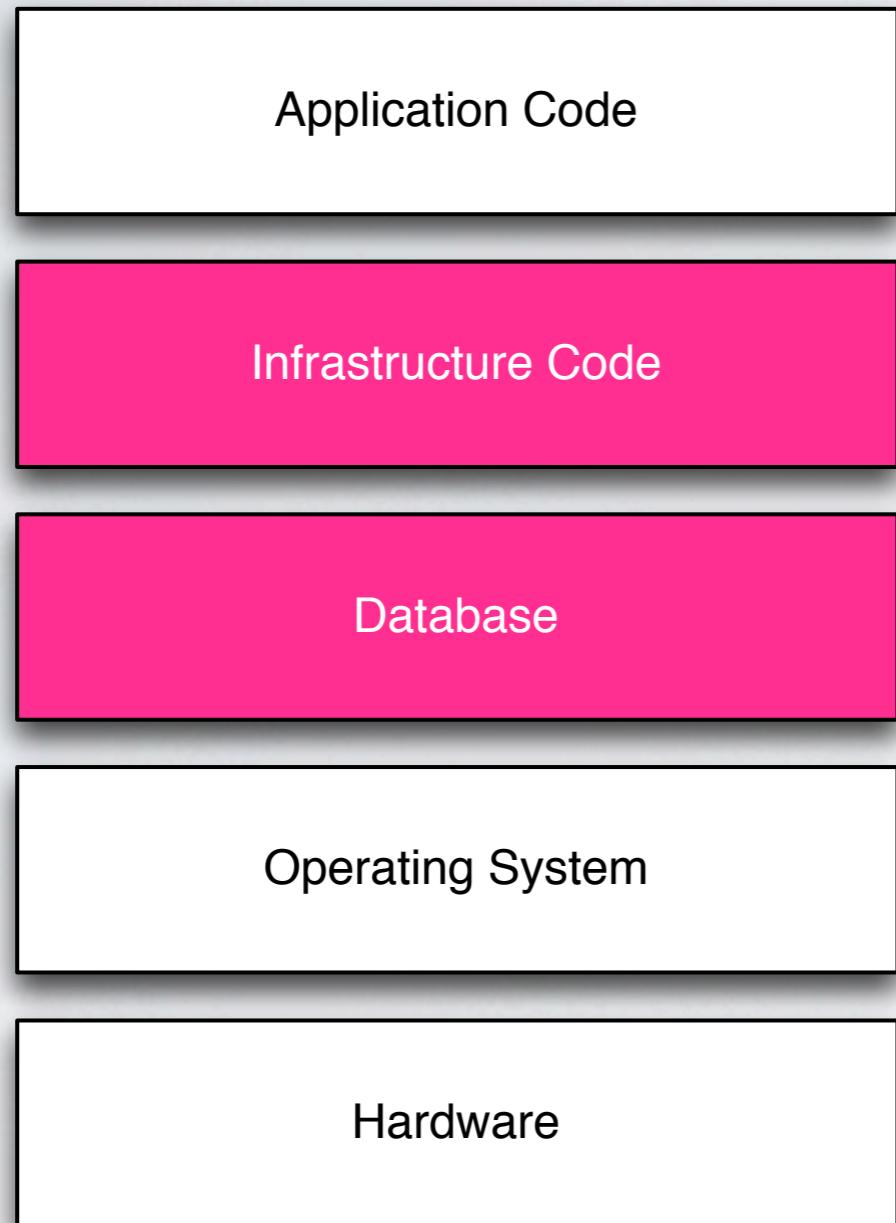


BUILDING A DATA PLATFORM WITH DATOMIC

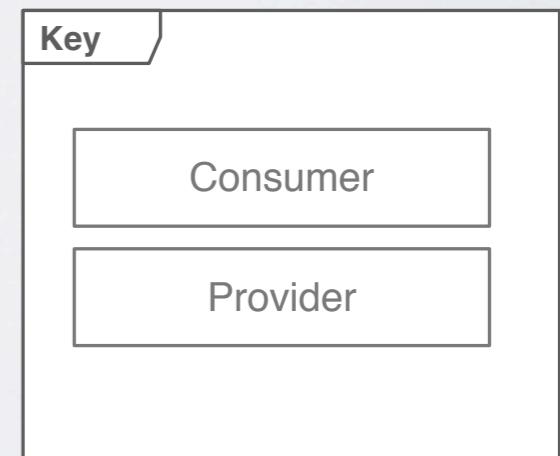
Antonio Andrade

Clojure/conj 2013

WHAT IS A DATA PLATFORM



Software system that defines a **data representation**, abstracts its **persistence** and **retrieval**, and provides tools to **manipulate** that data.



BACKGROUND

- Software Architect at Appian
- Build your own software
- Customers build their own applications and data models
- Highly dynamic system. Everything happens at runtime
- **Can't** make too **many modeling** decisions early on

WHERE DOES APPIAN POWER COME FROM?

- Transactional
- Rich data modeling: Arbitrary tree structures
- Data manipulation and query power: APL variant
- Data consistency: Multiple queries on data snapshot
- Performance. Ultra high-speed, in-memory, column-oriented data store
- Programming model: Embedded programming language

ROADMAP

- Data models & schemas
- Datomic data model
- A self-describing schema definition
- Techniques for versioning and data migration
- Enforcing data constraints
- Reified Transactions

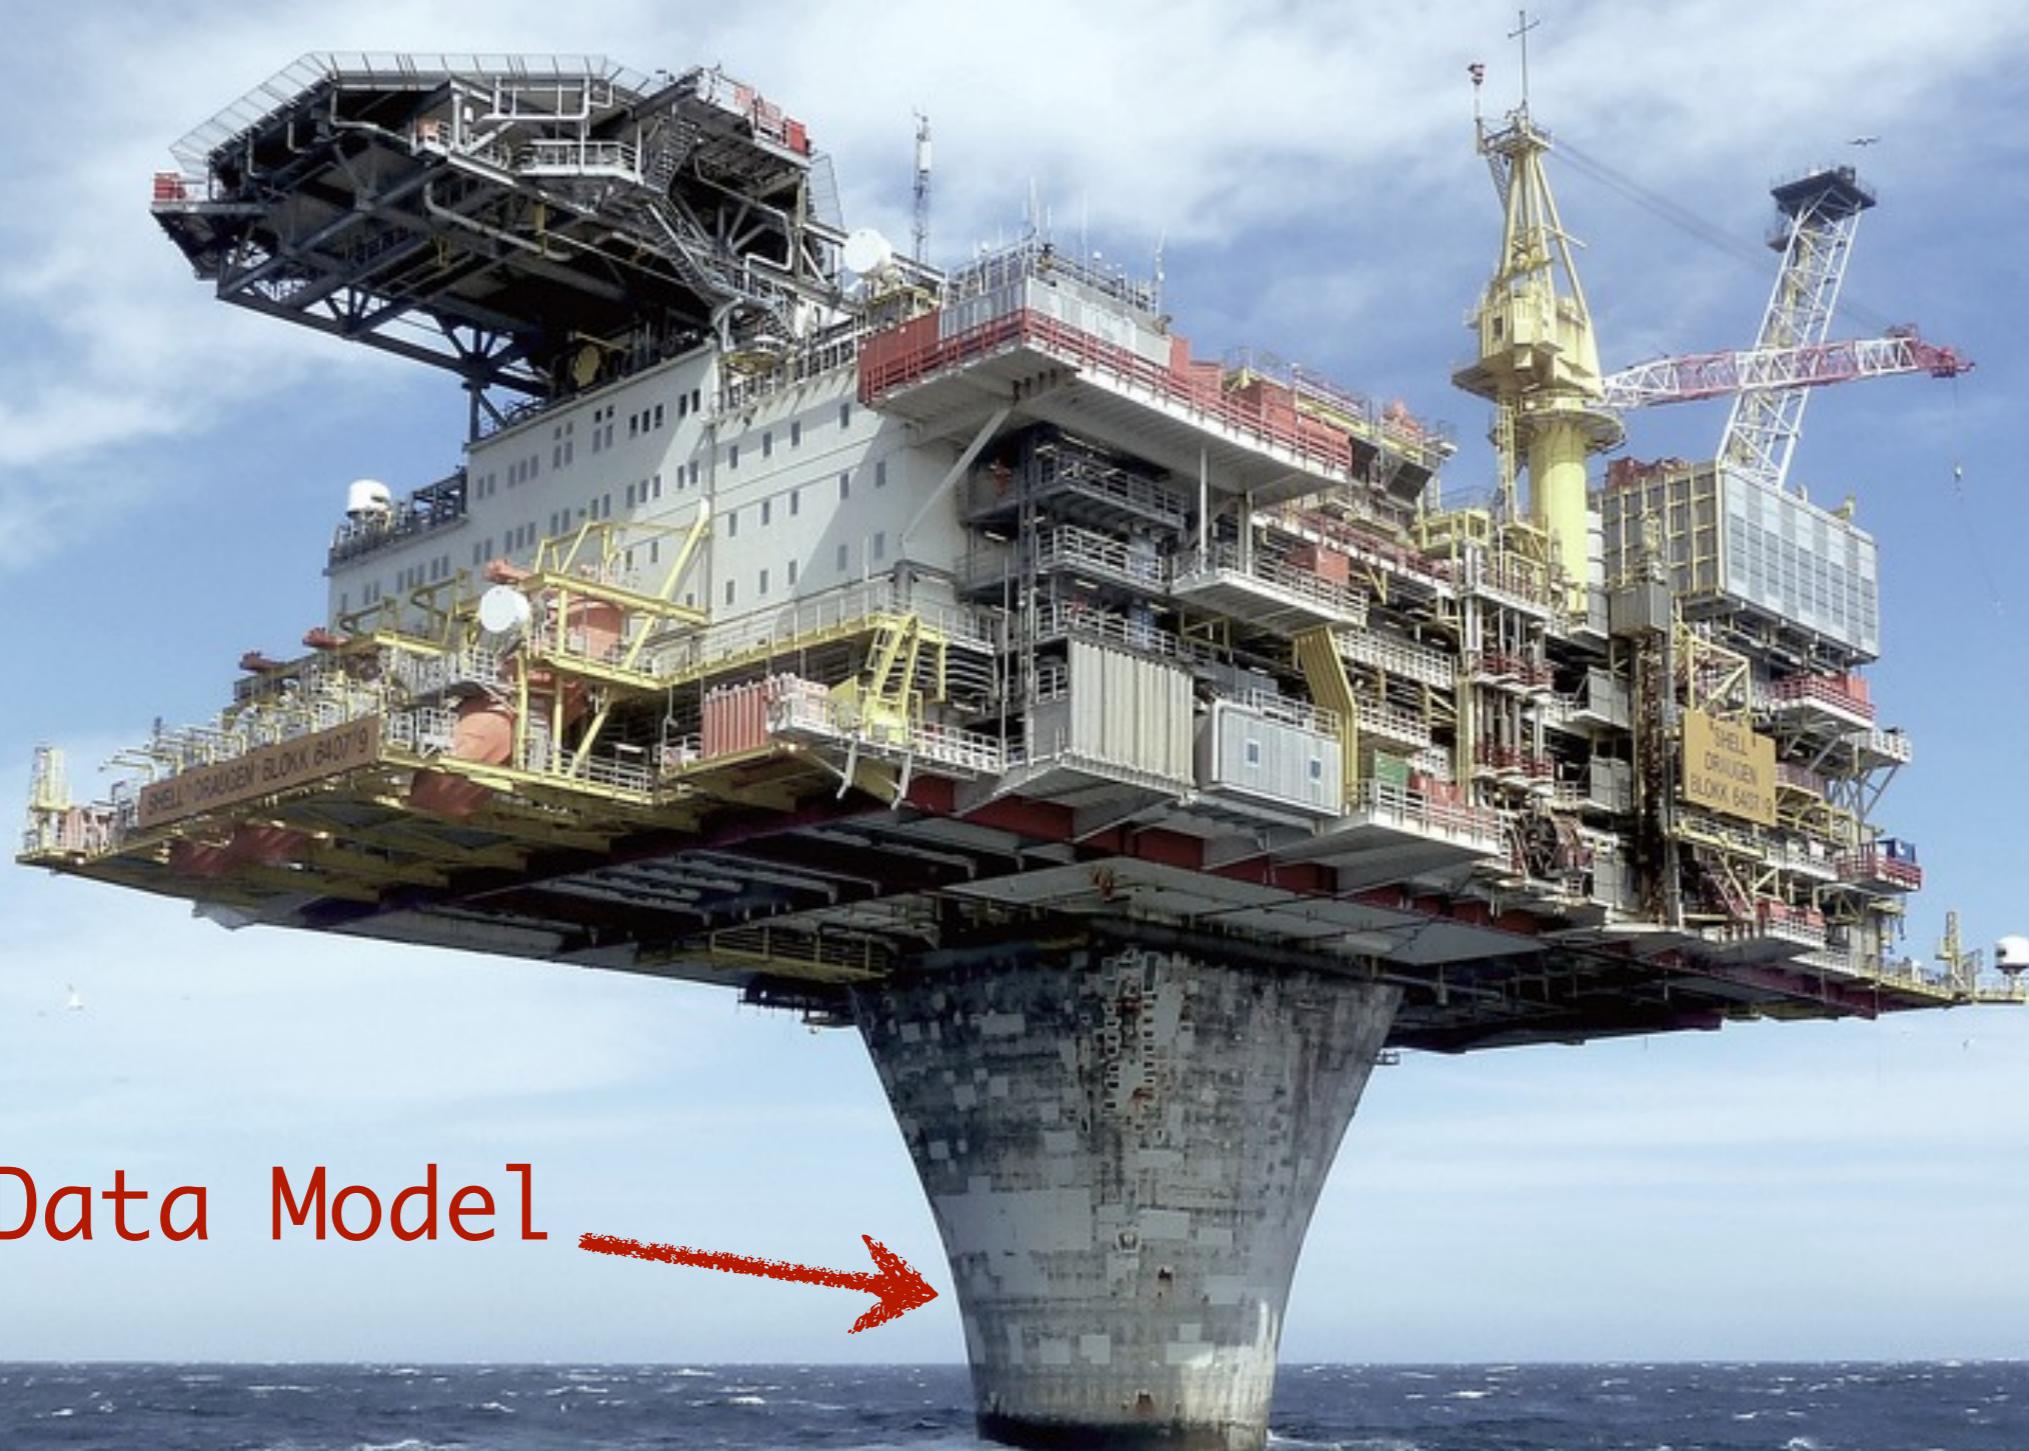
DATA MODELS

- Describe the **structure** of the *objects* represented by a system and the **relationship** between those *objects*.
- Always exists, either implicitly or explicitly.
- Really valuable as a communication tool and documentation.
- A schema is a *language* used to describe them.

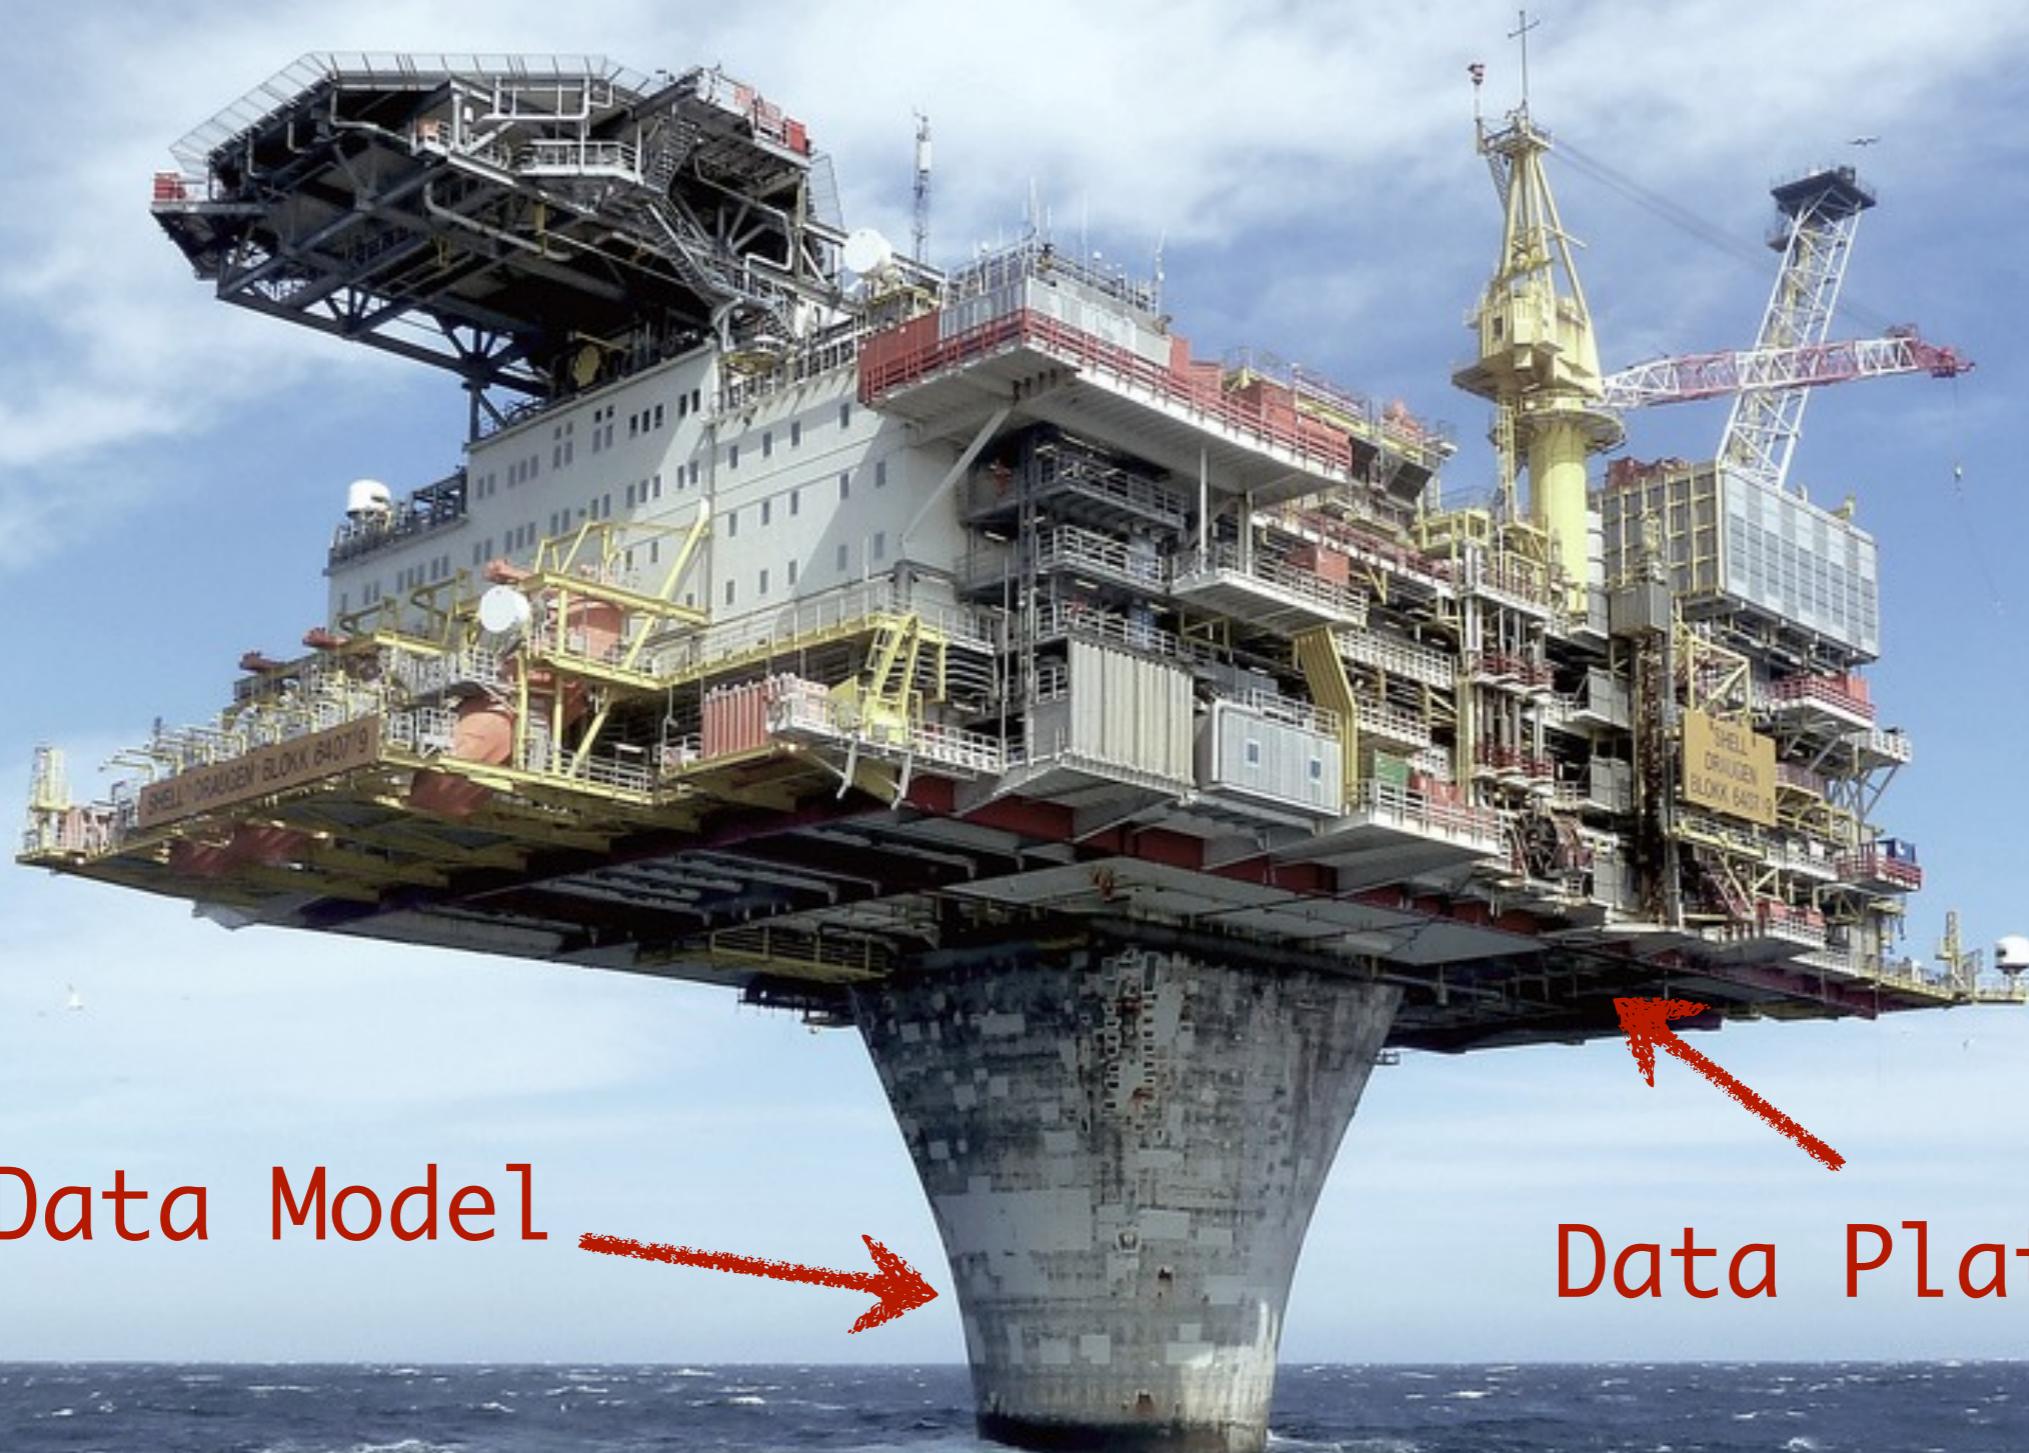
DATA MODELS ARE AT
THE CORE OF DATA
PLATFORMS



http://en.wikipedia.org/wiki/Draugen_oil_field



http://en.wikipedia.org/wiki/Draugen_oil_field

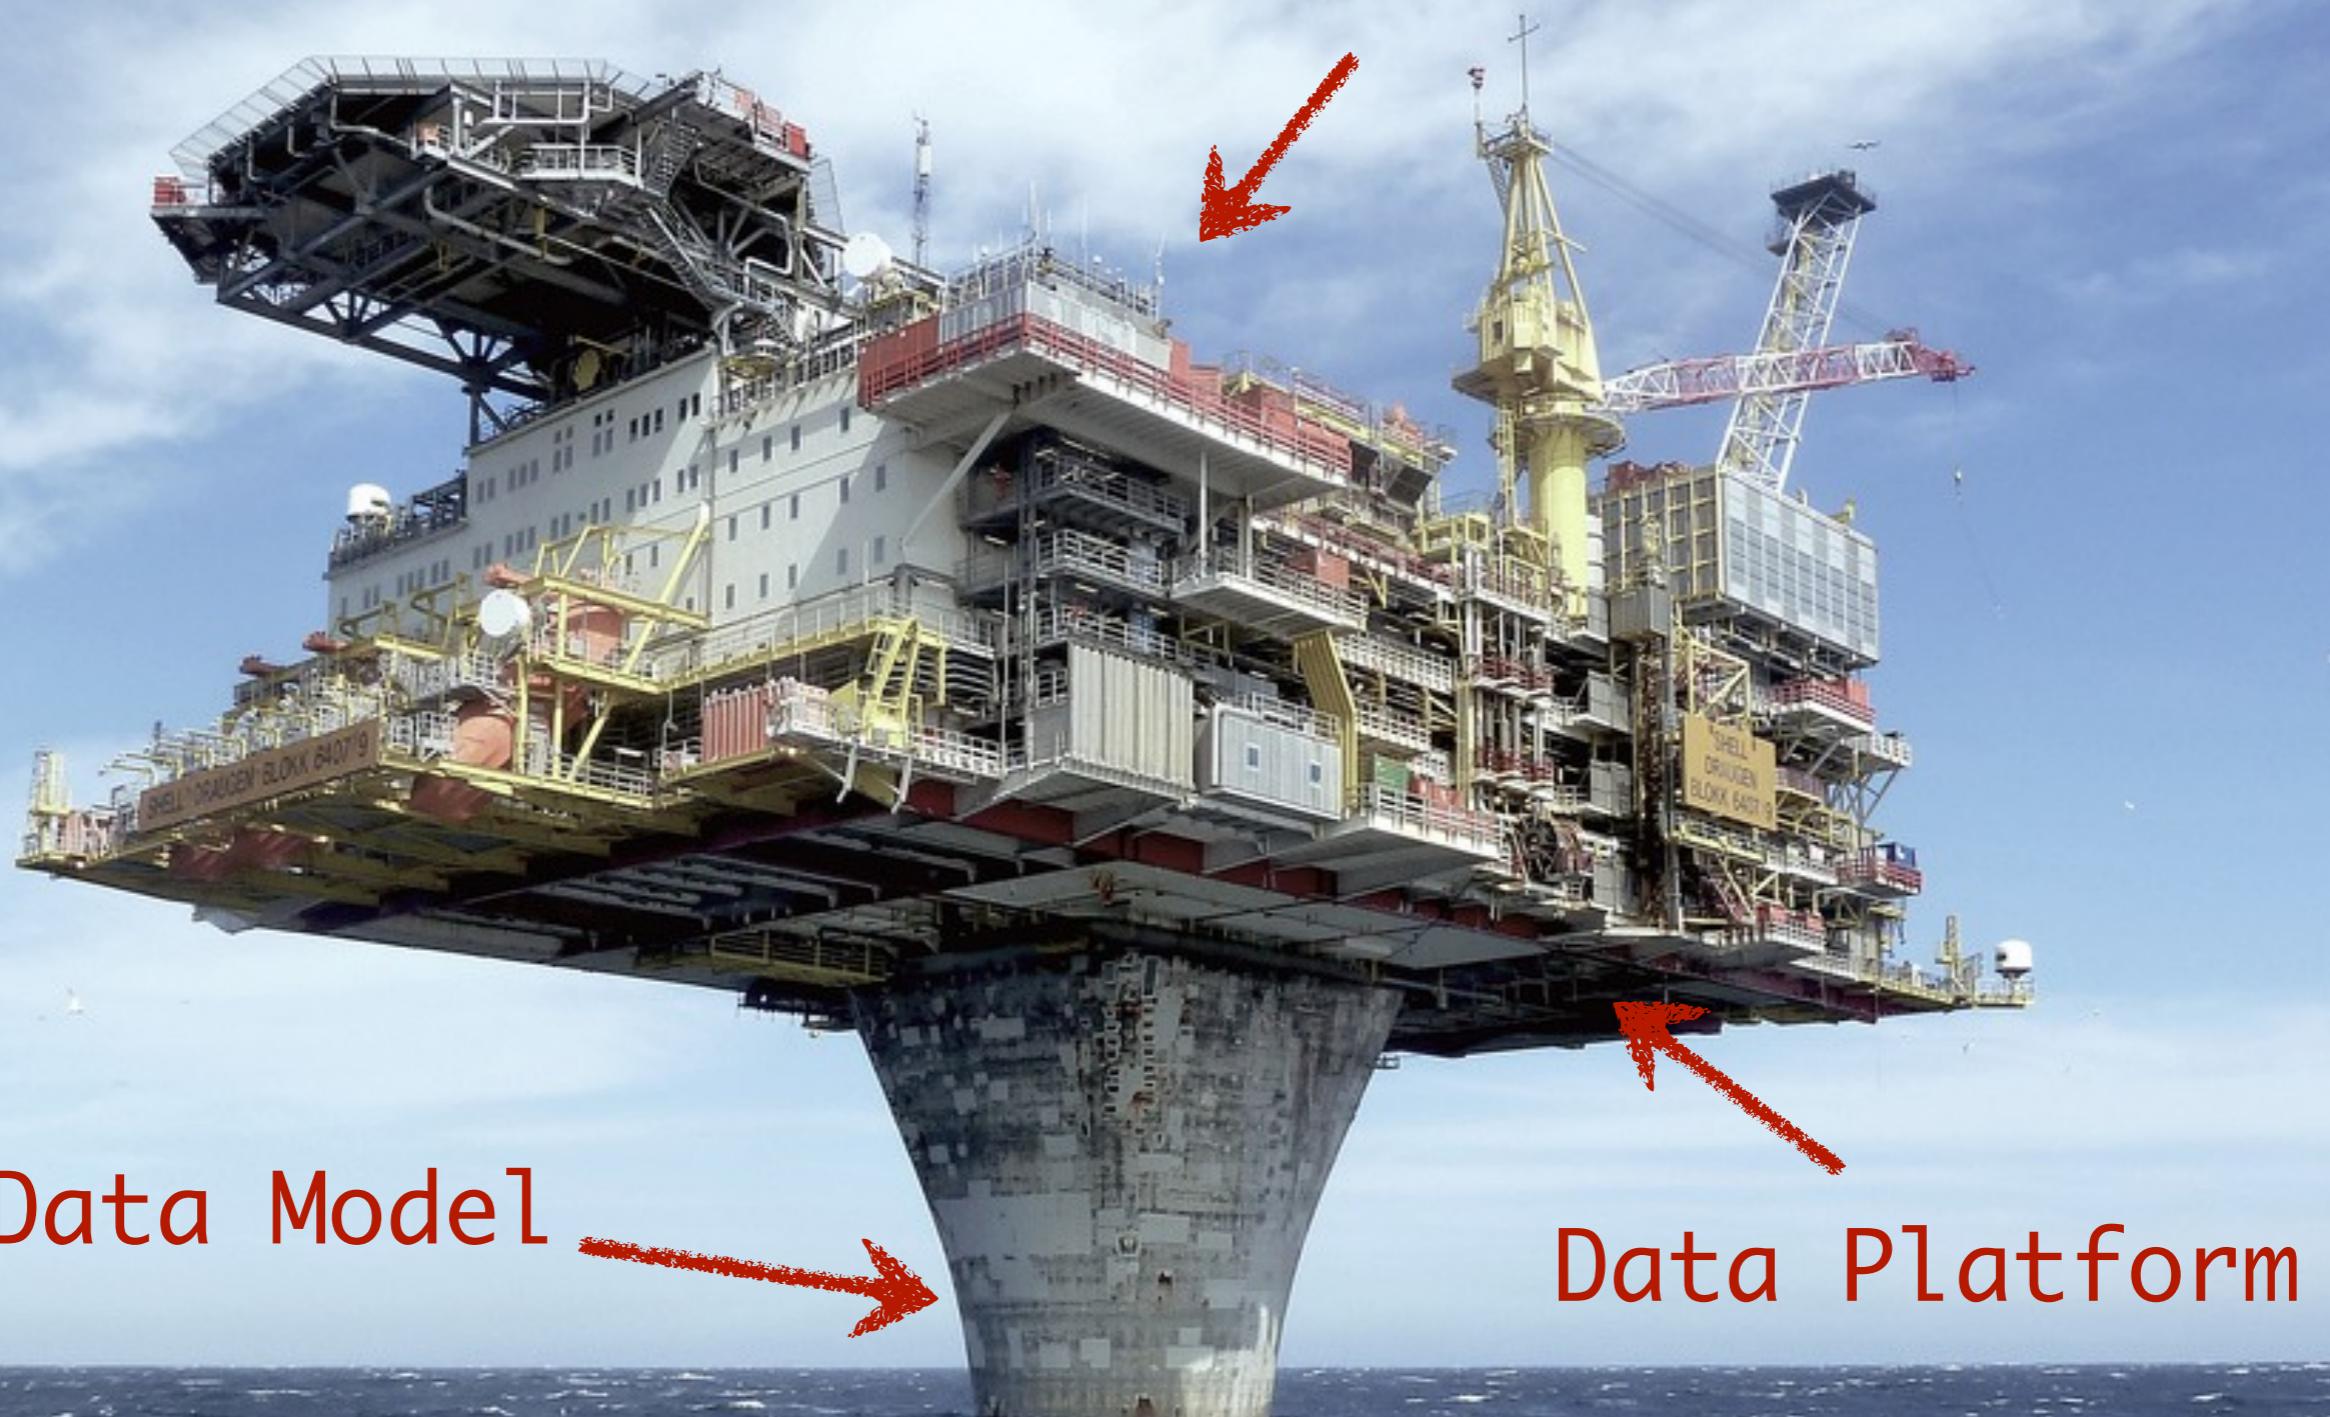


Data Model

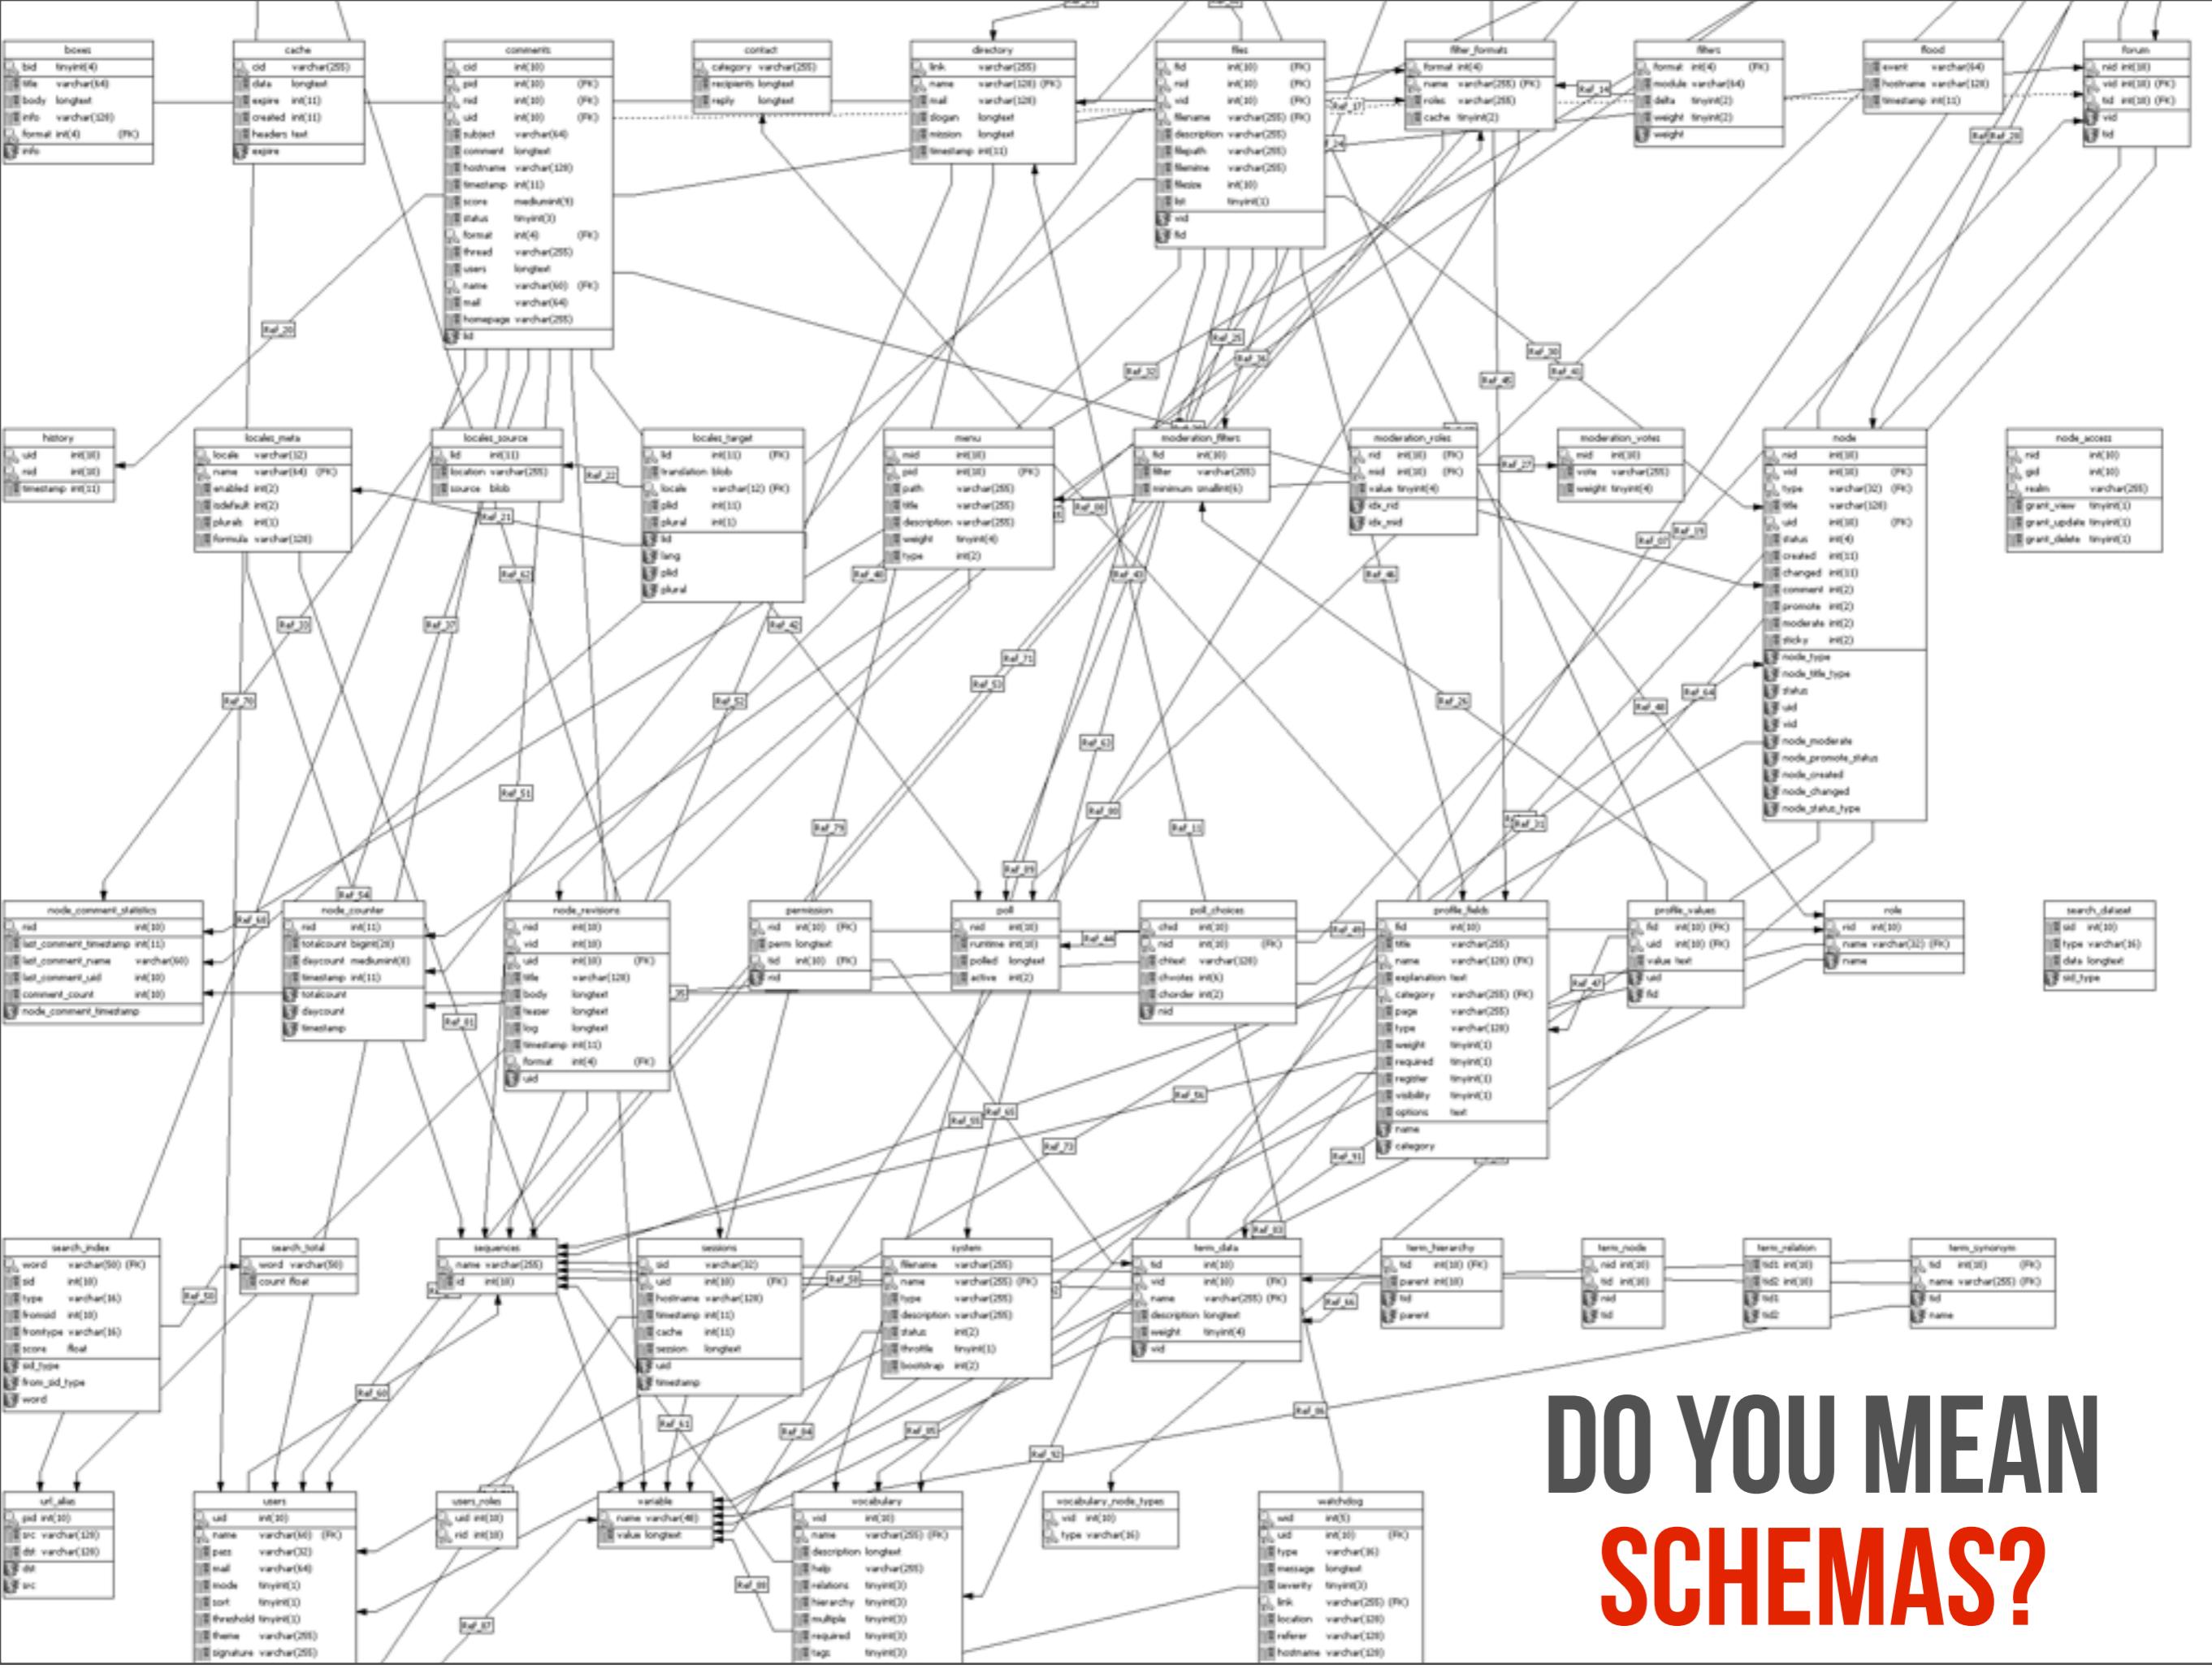
Data Platform

http://en.wikipedia.org/wiki/Draugen_oil_field

Application Code



http://en.wikipedia.org/wiki/Draugen_oil_field



DO YOU MEAN SCHEMAS?

SCHEMAS

- Good for performance
- Great for documentation. New hires appreciate it
- Good for enforcing data contracts
- Good as Reflection APIs on highly dynamic systems
- Good for code generation
- ***Really bad when structure needs to change***



Need to upgrade?

http://en.wikipedia.org/wiki/Draugen_oil_field

THE KEY IS ADAPTABILITY

- Application code should be able to retrieve the data in whatever shape it's needed.
- A system should not require live data migration to reshape the structure of the data when the structure needs to evolve.
- Avoid systems that can only serve efficient queries retrieving data in the shape they were originally stored in.

DATOM

THE DATOMIC DATA MODEL

BASIC API

;; require API and create alias
(require '[datomic.api :as d])

;; get a connection to the database
(def uri "datomic:mem://db-name")
(def conn (d/connect uri))

;; write (datoms) into the database
(d/transact conn
[DATOMS])

BASIC API

```
;; require API and create alias
(require '[datomic.api :as d])

;; get a connection to the database
(def uri "datomic:mem://db-name")
(def conn (d/connect uri))

;; write (datoms) into the database
(d/transact conn
  [DATOMS])
```

BASIC API

;; require API and create alias
`(require '[datomic.api :as d])`

;; get a connection to the database
`(def uri "datomic:mem://db-name")`
`(def conn (d/connect uri))`

;; write (datoms) into the database
`(d/transact conn`
`[DATOMS])`

DATOM == FACT

[*+/- entity attribute value tx*]

```
[[ :db/add #db/id[:db.part/user -1]
  :person/name "John"]
[ :db/add #db/id[:db.part/user -1]
  :person/age 25]]
```

DATOM == FACT

[+/- entity attribute value tx]

```
[[ :db/add #db/id[:db.part/user -1]
  :person/name "John"]
[ :db/add #db/id[:db.part/user -1]
  :person/age 25]]
```

DATOM == FACT

[+/- entity attribute value tx]

```
[[ :db/add #db/id[:db.part/user -1]
  :person/name "John"]
[ :db/add #db/id[:db.part/user -1]
  :person/age 25]]
```

DATOM == FACT

[+/- entity attribute **value** tx]

```
[[ :db/add #db/id[:db.part/user -1]
  :person/name "John"]
 [ :db/add #db/id[:db.part/user -1]
  :person/age 25]]
```

DATOM == FACT

[+/- entity attribute value tx]

```
[[ :db/add #db/id[:db.part/user -1]
  :person/name "John" ]
[ :db/add #db/id[:db.part/user -1]
  :person/age 25]
[:db/add #db/id[:db.part/tx]
  :db/txInstant (System/currentTimeMillis)]]
```

SHORT-HAND NOTATION

```
[{:db/id #db/id[:db.part/user]
  :person/name "John"
  :person/age 25}]
```

ATTRIBUTES ARE FACTS!

```
[{:db/id #db/id[:db.part/db]
  :db/ident :person/name
  :db/valueType :db.type/string
  :db/cardinality :db.cardinality/one
  :db.install/_attribute :db.part/db}]
```

Attributes are data and can be installed on
the fly

A
a
s
a
d
a
c
t

FREE ASSOCIATIONS

```
[{:db/id #db/id[:db.part/user]
  :person/name "John"
  :car/top-speed 200
  :computer/keyboard "DAS"
  :city/population 250000}]
```

What's this entity schema / data model?

HOW TO **DEFINE** A
HIGHER-LEVEL
SCHEMA AND USE IT
TO DEFINE **ITSELF**

CORE DATA MODEL: DATATYPE

ATTRIBUTE NAME	ATTRIBUTE TYPE	CARDINALITY
:dt/dt	ref	one
:dt/namespace	string	one
:dt/name	string	one
:db/parent	ref	one
:dt/list	ref	one
:dt/component	ref	one
:db/fields	ref	many

:dt/dt

```
[{:db/id #db/id[:db.part/db]
:db/ident :dt/dt
:db/valueType :db.type/ref
:db/cardinality :db.cardinality/one
:db/doc
```

“A reference to the data type of this entity. Entities with this attribute are known as **Typed-Entities**.¹”

```
:db.install/_attribute :db.part/db}]
```

≈ .getClass

:dt/namespace

```
[{:db/id #db/id[:db.part/db]
:db/ident :dt/namespace
:db/valueType :db.type/string
:db/cardinality :db.cardinality/one
:db/doc
  “Data types have qualified names.
  This is the data type namespace.”
:db.install/_attribute :db.part/db}]
```

```
≈ #(-> % .getPackage .getName)
```

:dt/name

```
[{:db/id #db/id[:db.part/db]
:db/ident :dt/name
:db/valueType :db.type/string
:db/cardinality :db.cardinality/one
:db/doc
  “The local part or unqualified name
  of a data type.”
:db.install/_attribute :db.part/db}]
```

≈ .getSimpleName

:dt/parent

```
[{:db/id #db/id[:db.part/db]
  :db/ident :dt/parent
  :db/valueType :db.type/ref
  :db/cardinality :db.cardinality/many
  :db/doc "A reference to the data type
that this data type inherits fields from."
  :db.install/_attribute :db.part/db}]
```

≈ .getSuperClass

:dt/list

```
[{:db/id #db/id[:db.part/db]
  :db/ident :dt/list
  :db/valueType :db.type/ref
  :db/cardinality :db.cardinality/one
  :db/doc "Dynamically-generated
reference to a list data type. Provides
support for multi-dimensional values."
  :db/install/_attribute :db.part/db}]
```

```
= #( -> (Array/newInstance % 0) .getClass)
```

:dt/component

```
[{:db/id #db/id[:db.part/db]
:db/ident :dt/component
:db/valueType :db.type/ref
:db/cardinality :db.cardinality/one
:db/doc Only populated if this is a
list data type."
:db.install/_attribute :db.part/db}]
```

≈ .getComponentType

:dt/fields

```
[{:db/id #db/id[:db.part/db]
:db/ident :dt/fields
:db/valueType :db.type/ref
:db/cardinality :db.cardinality/many
:db/doc
  “References to the data type fields,
  themselves data types.”
:db.install/_attribute :db.part/db}]
```

≈ .getFields

DATATYPE ≈ CLASS

ATTRIBUTE	IS ANALOGOUS TO
:dt/dt	.getClass
:dt/namespace	#(-> % .getPackage .getName)
:dt/name	.getSimpleName
:db/parent	.getSuperClass
:dt/list	#(-> (Array/newInstance % 0) .getClass)
:dt/component	.getComponentType
:db/fields	.getFields

PERSON INSTANCE

```
[{:db/id #db/id[:db.part/user]  
:person/name "John"  
:person/age 25}]
```

PERSON INSTANCE

```
[{:db/id #db/id[:db.part/user]  
 :dt/dt :person.datatype/entity  
 :person/name "John"  
 :person/age 25}]
```

A **Typed** Entity

dt/dt ≈ .getClass

Assume :person.datatype/entity had been defined

PERSON DATATYPE

```
[{:db/id :person.datatype/entity  
:dt/dt :dt/dt  
:dt/namespace "customer"  
:dt/name "Person"  
:dt/fields [:dt/dt  
             :person/name  
             :person/age]}]
```

Also a **Typed** Entity

dt/dt ≈ .getClass

DATATYPE DATATYPE

```
;; The dt/dt entity is both the definition
;; of a data type and an instance of itself
[{:db/id :dt/dt
  :dt/dt :dt/dt
  :dt/namespace "system"
  :dt/name "Datatype"
  :dt/fields [:dt/dt
    :dt/namespace :dt/name :dt/parent
    :dt/list :dt/component :dt/fields}]]
```

Also a **Typed** Entity

dt/dt ≈ .getClass

DATATYPE DATATYPE

```
;; The dt/dt entity is both the definition
;; of a data type and an instance of itself
[{:db/id :dt/dt
:dt/dt :dt/dt
:dt/namespace "system"
:dt/name "Datatype"
:dt/fields [:dt/dt
:dt/namespace :dt/name :dt/parent
:dt/list :dt/component :dt/fields}]]
```

DATATYPE DATATYPE

```
;; The dt/dt entity is both the definition
;; of a data type and an instance of itself
[{:db/id :dt/dt
  :dt/dt :dt/dt
  :dt/namespace "system"
  :dt/name "Datatype"
  :dt/fields [:dt/dt
    :dt/namespace :dt/name :dt/parent
    :dt/list :dt/component :dt/fields}]]
```

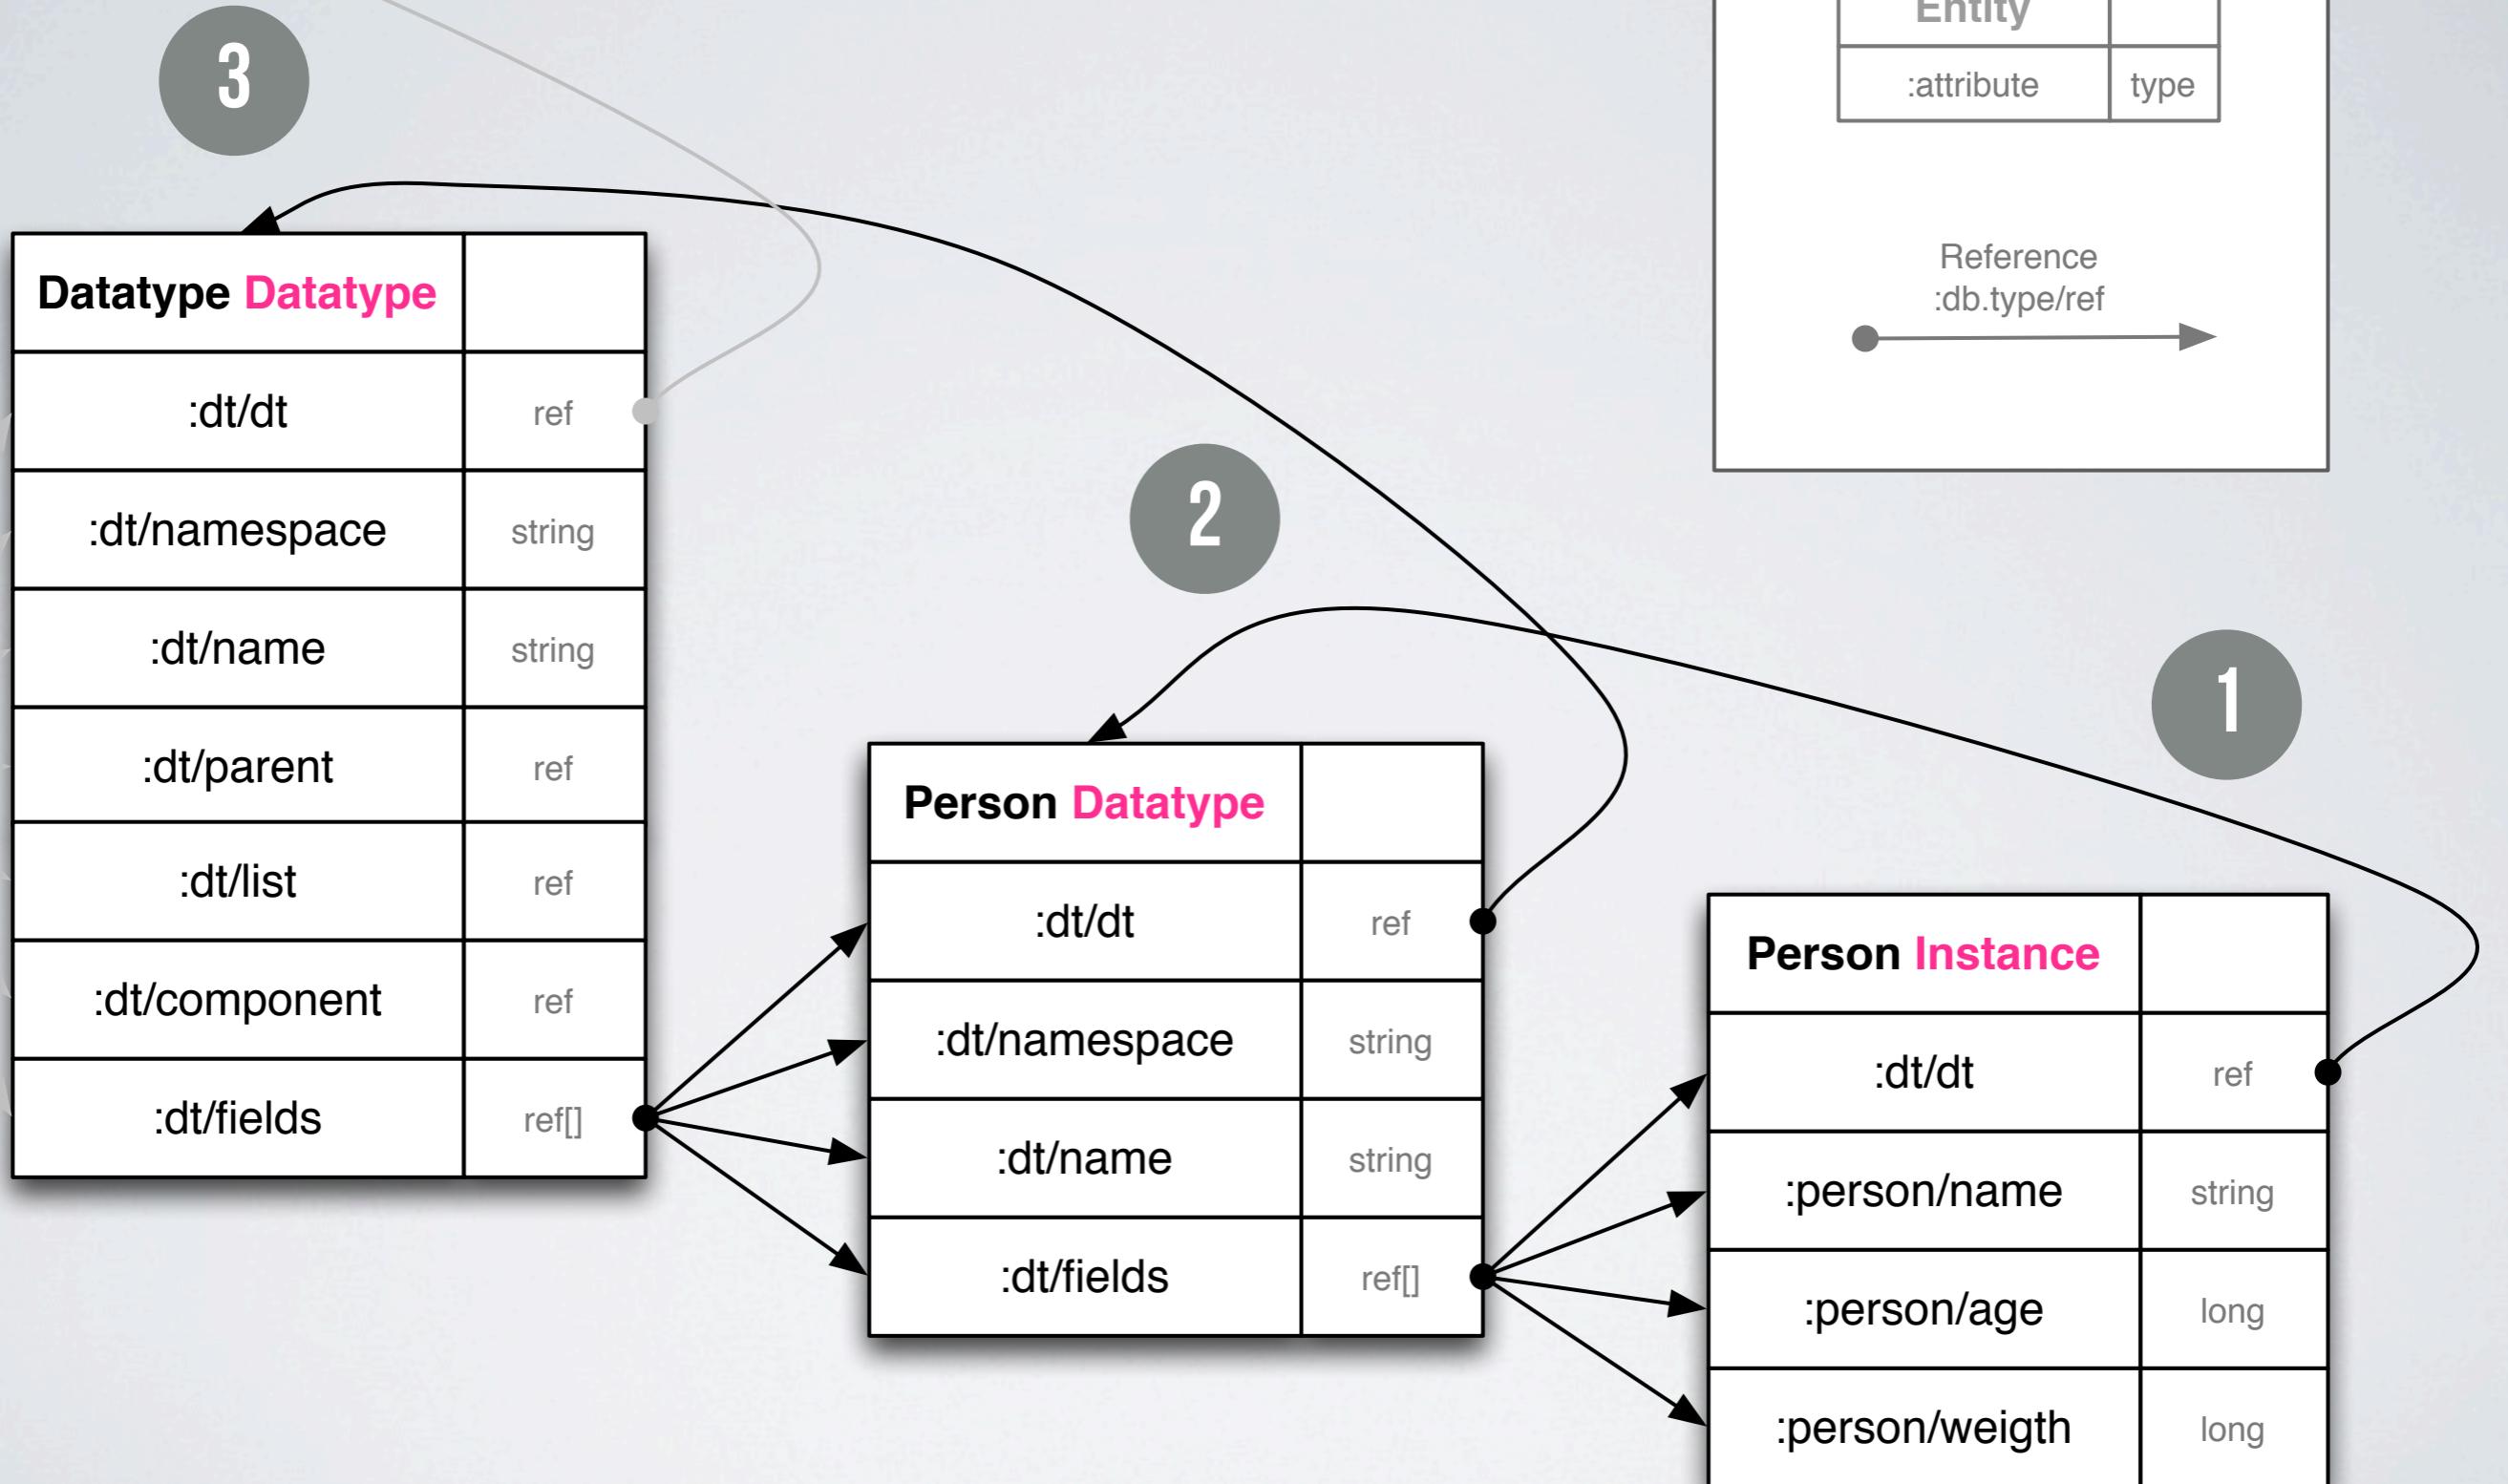
DATATYPE DATATYPE

```
;; The dt/dt entity is both the definition
;; of a data type and an instance of itself
[{:db/id :dt/dt
  :dt/dt :dt/dt
  :dt/namespace "system"
  :dt/name "Datatype"
  :dt/fields [:dt/dt
    :dt/namespace :dt/name :dt/parent
    :dt/list :dt/component :dt/fields}]]
```

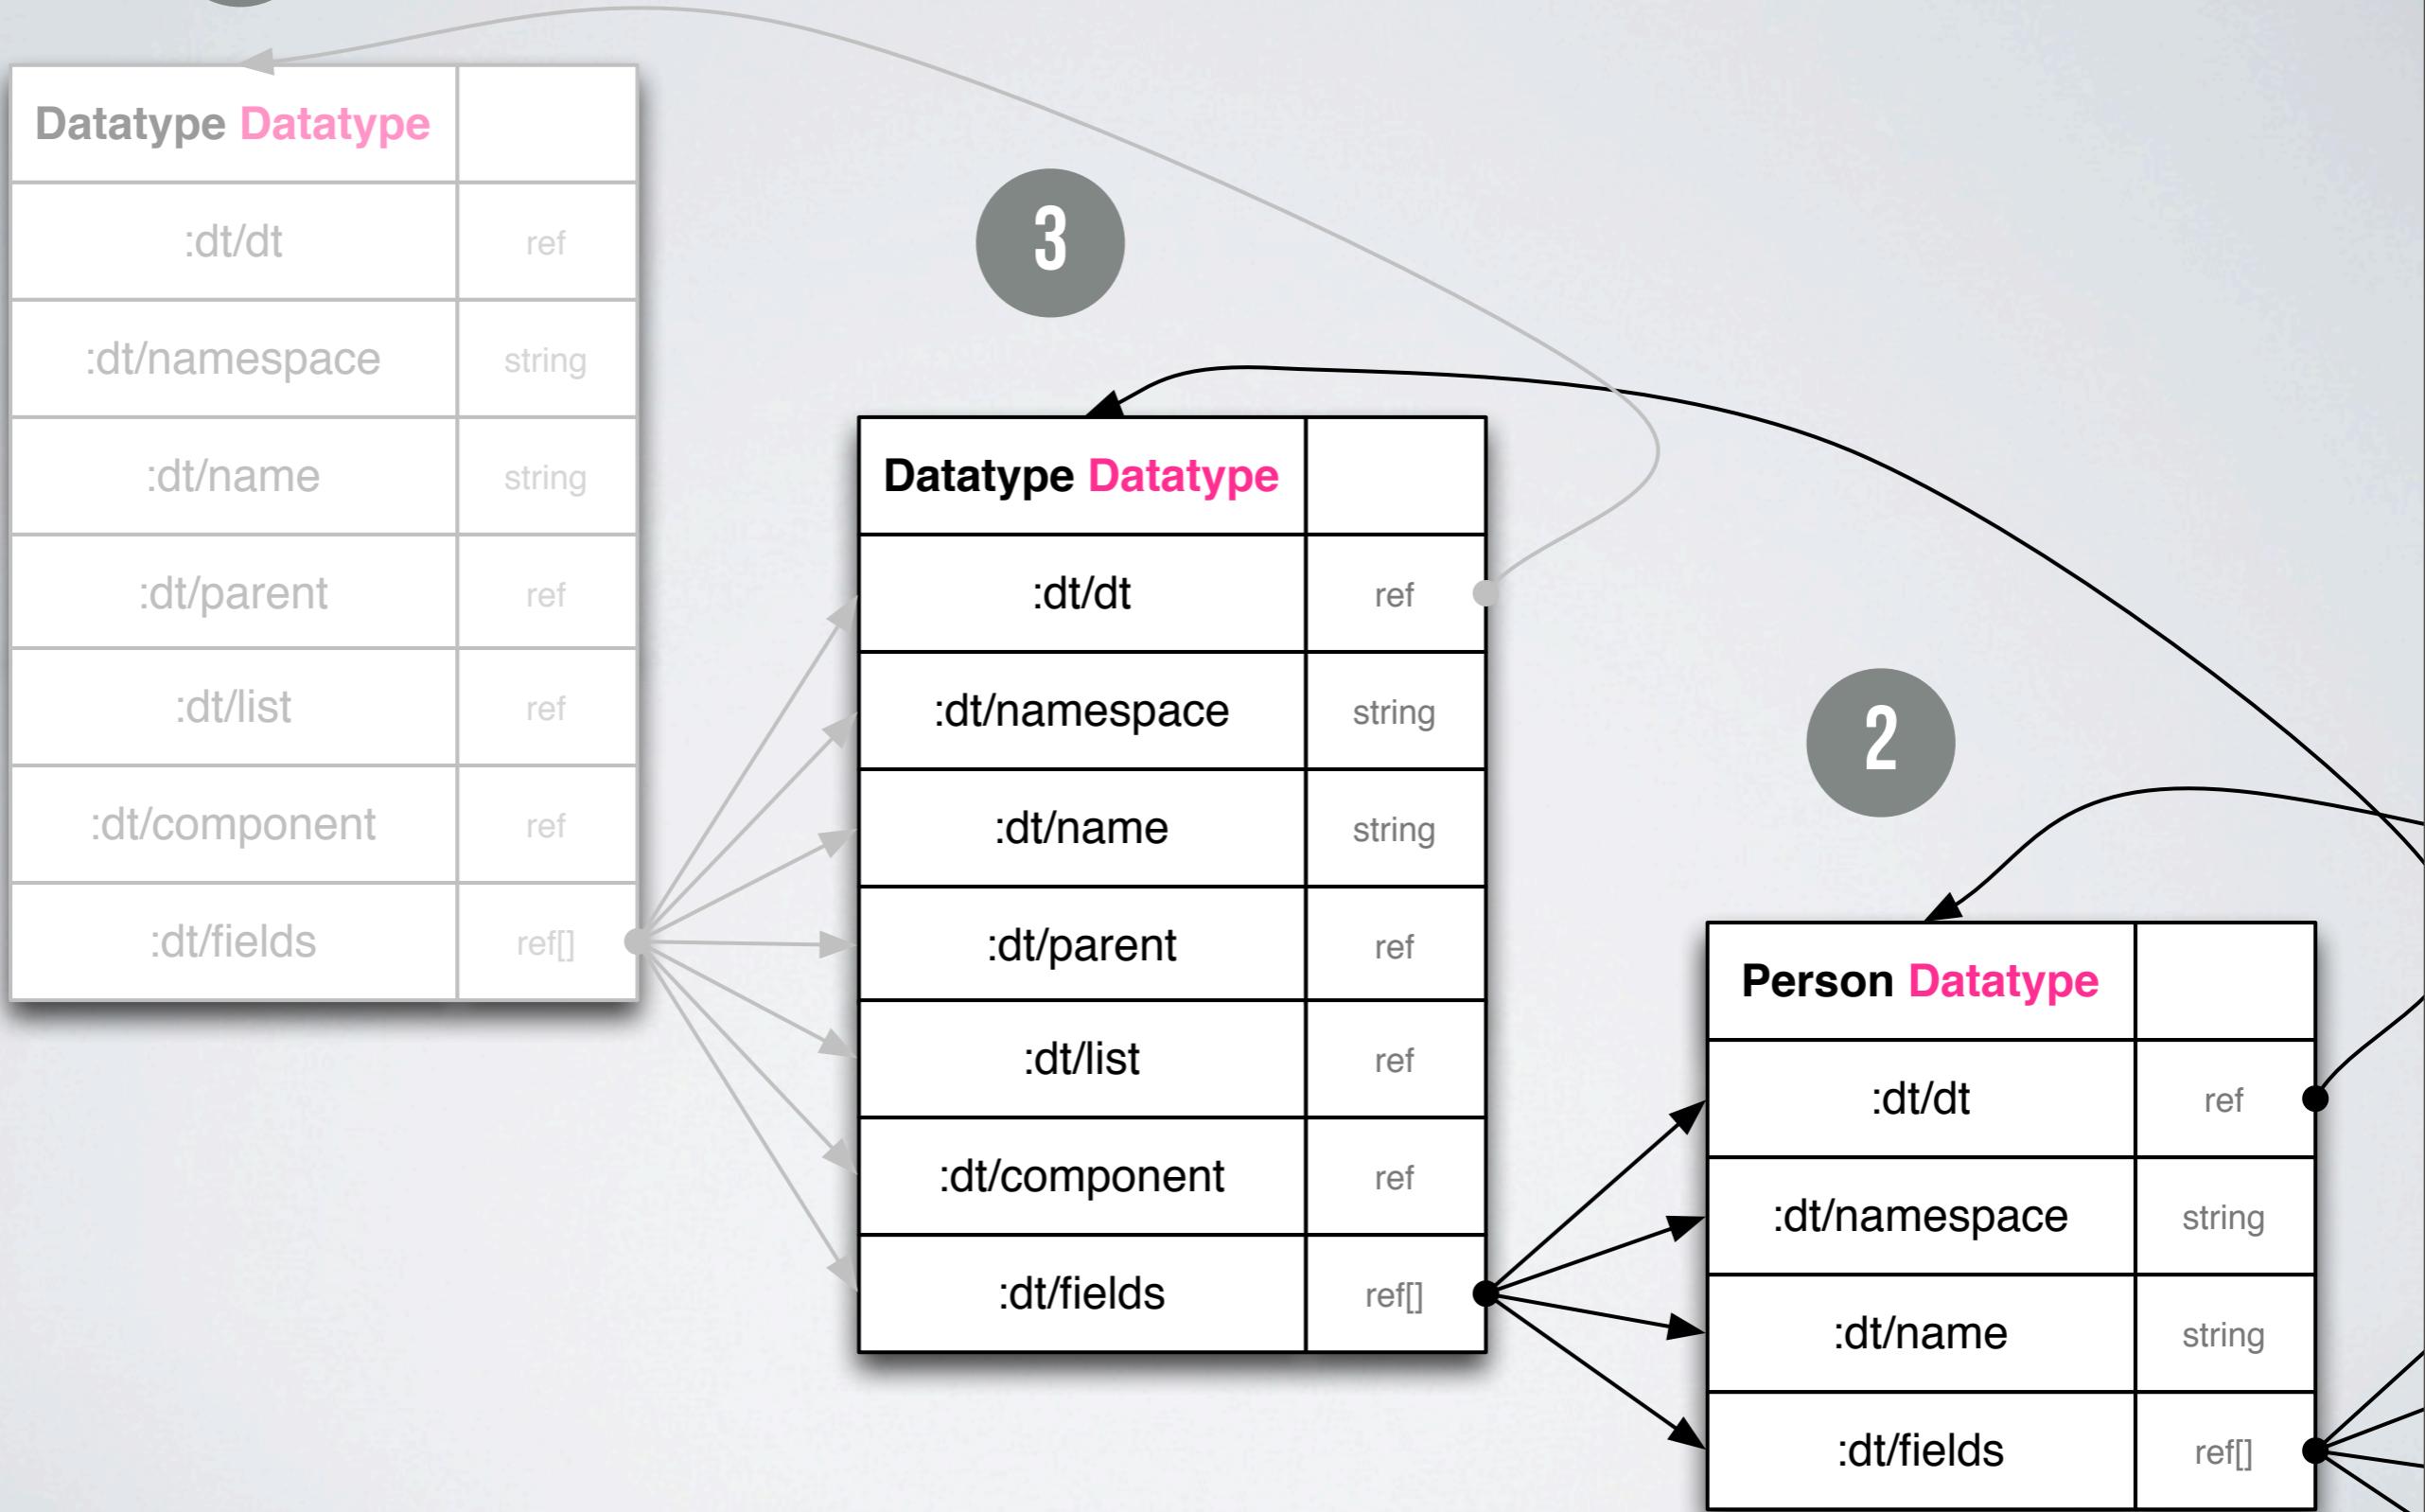
DATATYPE DATATYPE

```
;; The dt/dt entity is both the definition
;; of a data type and an instance of itself
[{:db/id :dt/dt
  :dt/dt :dt/dt
  :dt/namespace "system"
  :dt/name "Datatype"
  :dt/fields [:dt/dt
    :dt/namespace :dt/name :dt/parent
    :dt/list :dt/component :dt/fields}]]
```

META-MODEL



META-MODEL



DATATYPE HIERARCHIES

```
[{:db/id #db/id[:db.part/user]
:dt/dt :dt/dt
:dt/namespace "customer"
:dt/name "Student"
:dt/parent :person.datatype/entity
:dt/fields [:student/school}}]
```

:dt/parent ≈ .getSuperClass

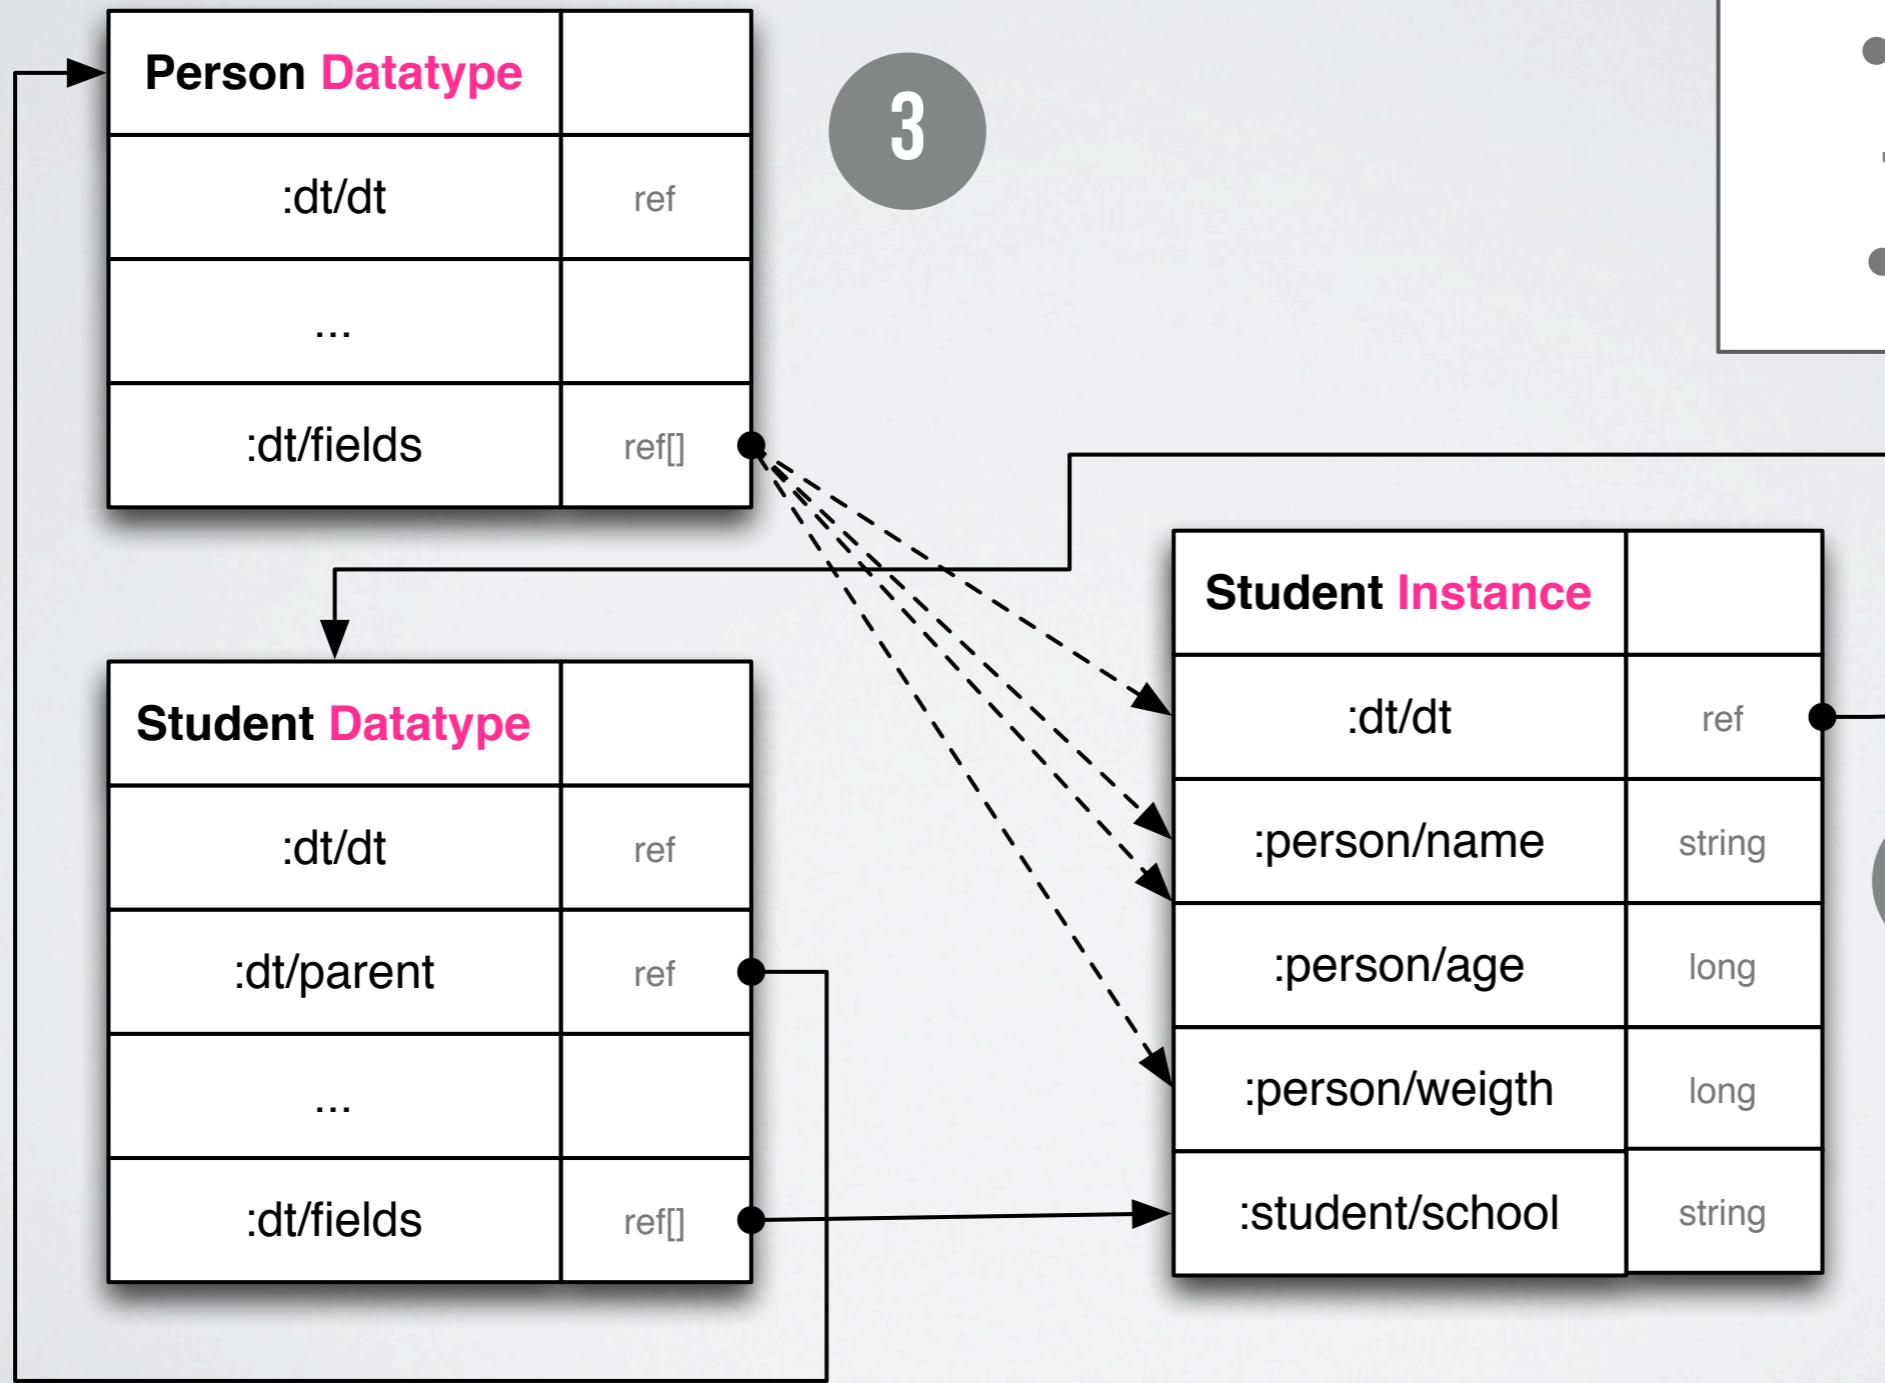
DATATYPE HIERARCHIES

```
[{:db/id #db/id[:db.part/user]
:dt/dt :dt/dt
:dt/namespace "customer"
:dt/name "Student"
:dt/parent :person.datatype/entity
:dt/fields [:student/school}}]
```

:dt/parent ≈ .getSuperClass

Assume *:student/school* had been defined

DATATYPE HIERARCHIES



MULTI-DIMENSIONAL VALUES

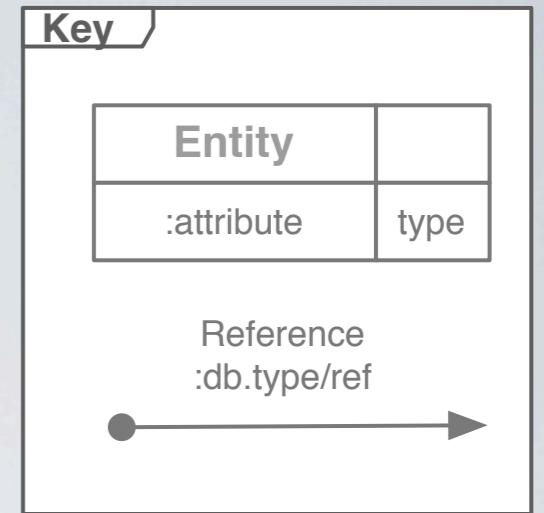
```
;; e.g. Person[]  
[{:db/id #db/id[:db.part/user -1]  
 :dt/dt :dt/dt  
 :dt/component :person.datatype/entity  
 :dt/_list :person.datatype/entity  
 :dt/fields [:dt.set/items]}  
  
;; e.g. Person[][]  
{:db/id #db/id [:db.part/user]  
 :dt/dt :dt/dt  
 :dt/component #db/id[:db.part/user -1]  
 :dt/_list #db/id[:db.part/user -1]  
 :dt/fields [:dt.set/items]}]
```

All references in Datomic are bidirectional so these refs are, in fact, redundant.

MULTI-DIMENSIONAL VALUES

```
;; e.g. Person[]  
[{:db/id #db/id[:db.part/user -1]  
 :dt/dt :dt/dt  
 :dt/component :person.datatype/entity  
 :dt/_list :person.datatype/entity  
 :dt/fields [:dt.set/items]}  
;; e.g. Person[][]  
{:db/id #db/id [:db.part/user]  
 :dt/dt :dt/dt  
 :dt/component #db/id[:db.part/user -1]  
 :dt/_list #db/id[:db.part/user -1]  
 :dt/fields [:dt.set/items]}]
```

You can use temporary ids to link entities before they are transacted.



3

Person[][] Datatype	
:dt/dt	ref
...	

2

Person[] Datatype	
:dt/dt	ref
:dt/list	
...	

1

Person Datatype	
:dt/dt	ref
:dt/list	ref
...	
:dt/fields	ref[]

Person[][] Instance	
:dt/dt	ref
:dt.set/items	ref[]

6

Person[] Instance	
:dt/dt	ref
:dt.set/items	ref[]

5

Person Instance	
:dt/dt	ref
:person/name	string
...	

4

All multi-dimensional entities
have a `:dt.set/items`
attribute

ANY TYPE / VARIANT TYPE

```
[{:db/id #db/id[:db.part/db]
  :db/ident :dt/any
  :db/valueType :db.type/ref
  :db/cardinality :db.cardinality/one
  :db/doc "Built-in types enumeration"
  :db.install/_attribute :db.part/db}
{:db/id #db/id[:db.part/db]
  :db/ident :dt.any/string
  :db/valueType :db.type/string
  :db/cardinality :db.cardinality/one
  :db.install/_attribute :db.part/db}
...
]
```

ANY TYPE / VARIANT TYPE

```
[{:db/id :dt/any
  :dt/dt :dt/dt
  :dt/namespace "system"
  :dt/name "Any"
  :dt/fields [:dt/dt
    :dt/any
    :dt.any/string
    :dt.any/boolean
    :dt.any/long
    ...
    :dt.any/ref}]]
```

Only one of the `dt.any/` attributes is present at runtime for any given “Any” entity*

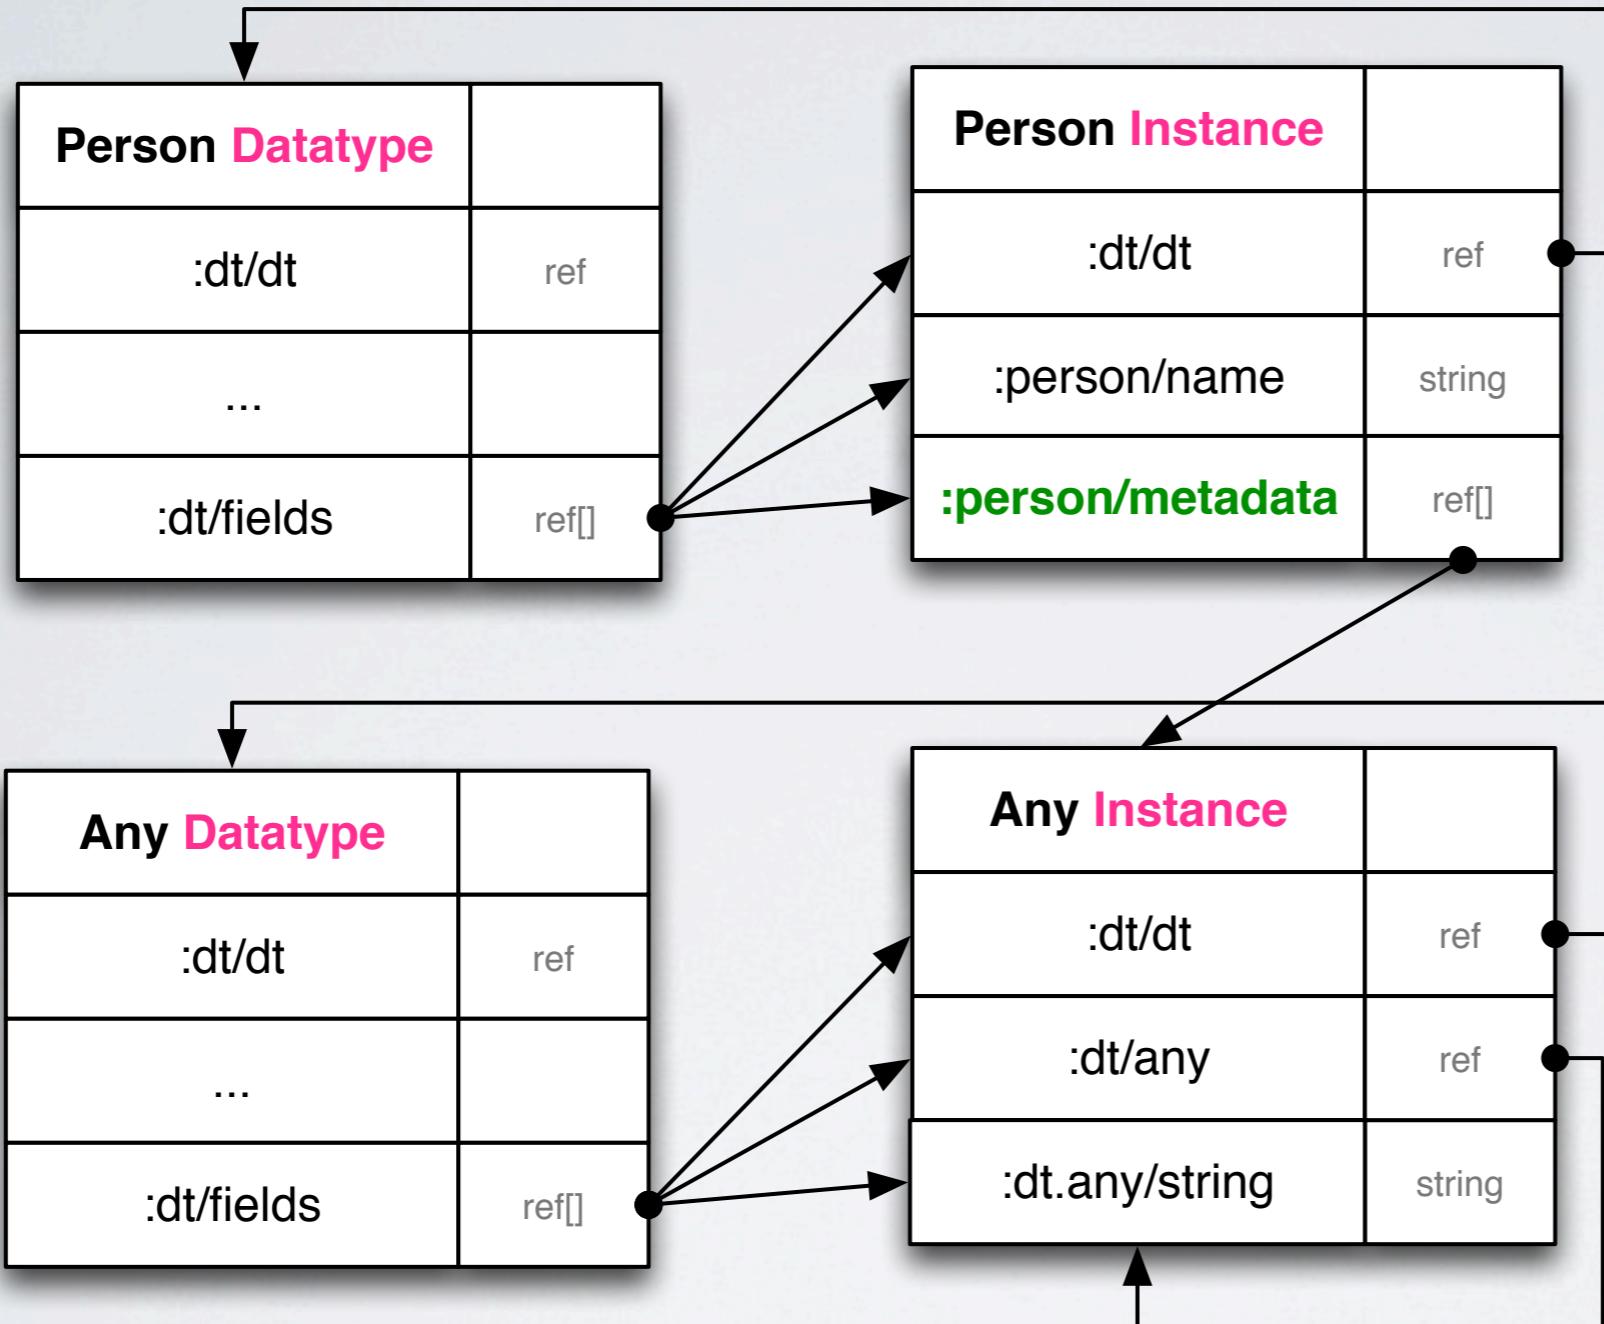
TYPED REFERENCES

```
[{:db/id #db/id[:db.part/db]
  :db/ident :person/metadata
  :db/valueType :db.type/ref
  :dt.meta/valueType :dt/any
  :db/cardinality :db.cardinality/many
  :db.install/_attribute :db.part/db}]
```

```
[{:db/id :person.datatype/entity
  :dt/fields [:person/metadata]}]
```

Assume :dt.meta/valueType had been defined

ANY TYPE / VARIANT TYPE



Only one of the *dt.any/** attributes is present at runtime for any given “Any” entity

HOW TO DEAL WITH CHANGES IN SCHEMAS

PERSON DATATYPE v2

```
[{:db/id :person.datatype/entity  
:dt/fields [:person/weight]}]
```

Transactions can add individual values for multi-valued attributes.

EXISTING PERSON INSTANCE

```
[{:db/id #db/id[:db.part/user]  
:dt/dt :person.datatype/entity  
:person/name "John"  
:person/age 25}]
```



all existing entities
are still valid!

PERSON DATATYPE v3

```
[{:db/id #db/id[:db.part/db]
:db/ident :health/insurance
:db/valueType :db.type/string
:db/cardinality :db.cardinality/one
:dt.meta/required true
:db.install/_attribute :db.part/db}]
```

Assume :dt.meta/required had been defined

PERSON DATATYPE v3

```
[{:db/id :person.datatype/entity  
:dt/fields [:health/insurance]}]
```



health insurance is NOT optional

EXISTING PERSON INSTANCE

```
[{:db/id #db/id[:db.part/user]  
:dt/dt :person.datatype/entity  
:person/name "John"  
:person/age 25}]
```



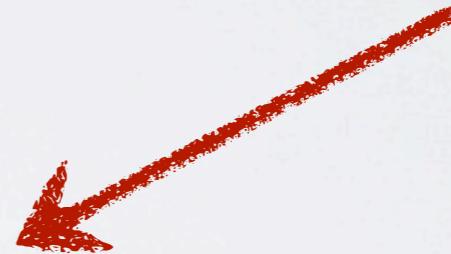
existing entities are
NO LONGER “valid”!

HOW TO FIX IT?

DATABASE FUNCTIONS

```
[{:db/id :health/insurance  
:dt.meta/migration  
(d/function  
'{:lang :clojure  
:params [db entity]  
:code [[  
:db/add entity  
:health/insurance "PPACA"]]]})}]
```

Must eventually return
a list of Datoms



Database functions!

DATABASE FUNCTIONS

```
[{:db/id #db/id[:db.part/db]
:db/ident :dt.meta/migration
:db/valueType :db.type/fn
:db/cardinality :db.cardinality/one
:db.install/_attribute :db.part/db}]
```

Datomic first-class functions

2

DECIDE WHEN TO VALIDATE

2

DECIDE WHEN TO VALIDATE

- Enforce that all required attributes include a migration function at the time the attributes is created.

DECIDE WHEN TO VALIDATE

- Enforce that all required attributes include a migration function at the time the attributes is created.
- Enforce constraints when **new** entities are created.

DECIDE WHEN TO VALIDATE

- Enforce that all required attributes include a migration function at the time the attributes is created.
- Enforce constraints when **new** entities are created.
- How? wrap (d/transact) and perform these checks.

3

DECIDE WHEN TO MIGRATE

3

DECIDE WHEN TO MIGRATE

- At query time, **but** cannot be used as a rule.

3

DECIDE WHEN TO MIGRATE

- At query time, **but** cannot be used as a rule.
- At transaction time, **but** queries could return “invalid” data.

3

DECIDE WHEN TO MIGRATE

- At query time, **but** cannot be used as a rule.
- At transaction time, **but** queries could return “invalid” data.
- At attribute access time, **but** restricted scope.

DECIDE WHEN TO MIGRATE

- At query time, **but** cannot be used as a rule.
- At transaction time, **but** queries could return “invalid” data.
- At attribute access time, **but** restricted scope.
- As a background processing task, **but** data would be eventually consistent.

DECIDE WHEN TO MIGRATE

- At query time, **but** cannot be used as a rule.
- At transaction time, **but** queries could return “invalid” data.
- At attribute access time, **but** restricted scope.
- As a background processing task, **but** data would be eventually consistent.
- Manually, driven by application code.

HOW TO ENFORCE ARBITRARY DATA CONSTRAINTS

ENUMS

```
[{:db/id #db/id[:db.part/db]
  :db/ident :person/gender
  :db/valueType :db.type/ref
  :db/cardinality :db.cardinality/one
  :db.install/_attribute :db.part/db}]
```

```
{:db/id #db/id[:db.part/user]
  :db/ident :person.gender/female}
```

```
{:db/id #db/id[:db.part/user]
  :db/ident :person.gender/male}]]
```

ENUMS

```
[[:db/add #db/id[:db.part/user]  
  :person/gender :car/top-speed]]
```



will succeed

HOW TO FIX IT?

```
[{:db/id #db/id [:db.part/user]
  :db/ident :add-fk
  :db/fn (d/function
    '{:lang :clojure
      :params [db e a v]
      :code (let [;; get value's symbolic keyword
                 ident (if (keyword? v) v (d/ident db v))
                 ;; get enum's constraint
                 aent (d/entity db a)
                 fname (:db/ident aent)
                 enum-ns (:enum/ns aent) ;;; find all possible enum idents
                 allowed (if (nil? enum-ns)
                           (throw (Exception. (str "Cannot check fk
constraint. " fname " has no :enum/ns attribute"))))
                           (d/q '[:find ?ident
                                 :in $ enum-ns
                                 :where
                                 [_ :db/ident ?ident]
                                 [(namespace ?ident) ?ns]
                                 [(= ?enum-ns ?ns)]]
                                 db enum-ns))]
                  ;; enforce constraint
                  (if (contains? allowed ident)
                      [[:db/add e a v]]
                      (throw (Exception. (str v " is not one of " allowed))))))}]
```

1 give it a name

2 add hint to enum values

3 look up valid enum values

4 fail transaction

ENUMS

```
[[:db/add-fk #db/id[:db.part/user]  
:person/gender :car/top-speed]]
```



will FAIL

REIFIED TRANSACTIONS

AUDIT TRAIL?

[[db/retract 12345 :health/insurance]]

TRANSACTIONS ARE FACTS!

```
[[db/retract 12345 :health/insurance  
[ :db/add #db/id[:db.part/tx]  
  :audit/source username-id] ]
```

Assume :audit/source had been defined

THANK YOU

Antonio Andrade

Clojure/conj 2013